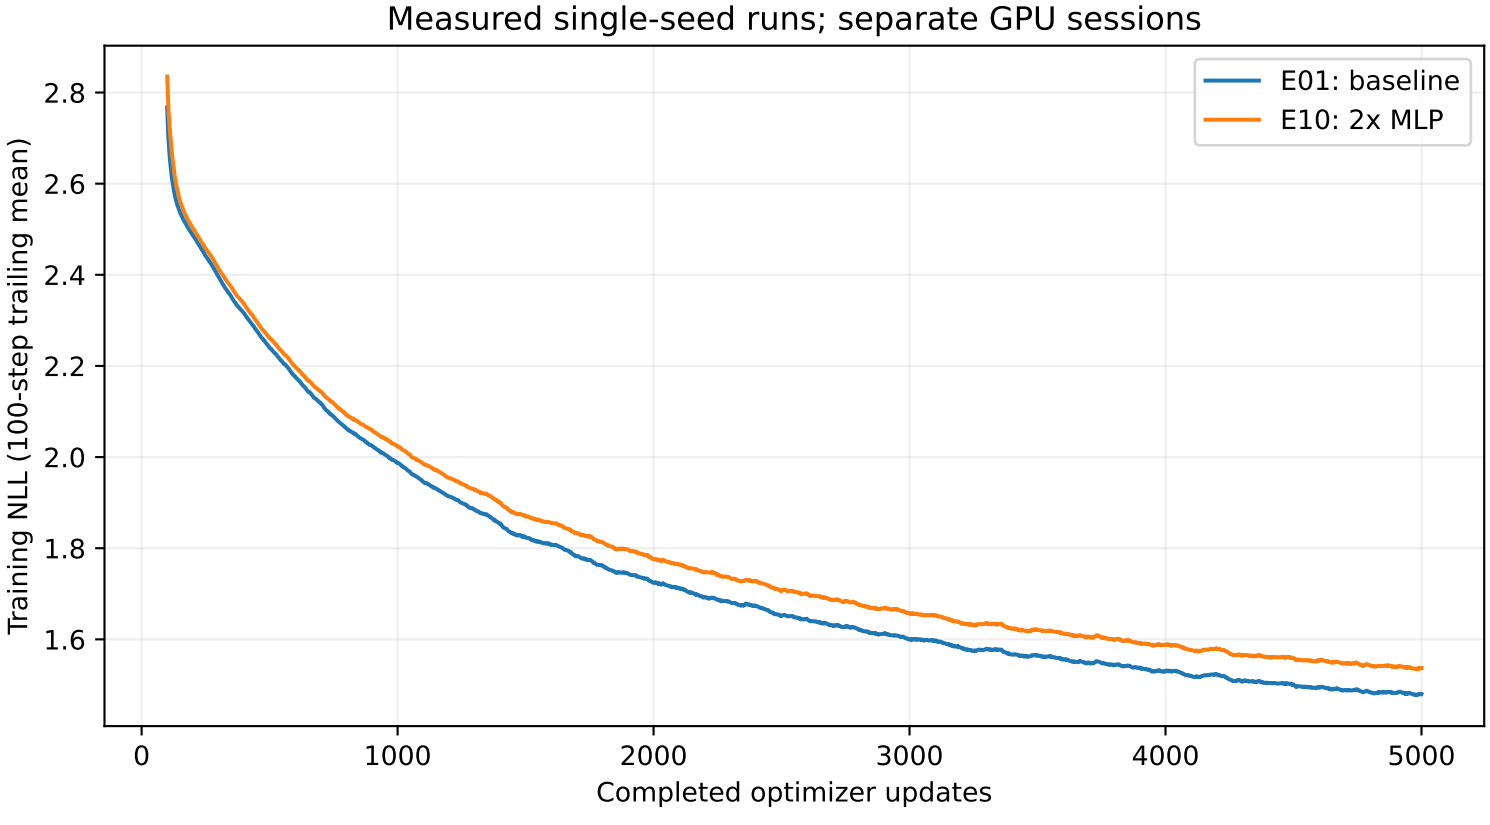
Measured single-seed runs; separate GPU sessions
E01: baseline
E10: 2x MLP
Training NLL (100-step trailing mean)
2.8
2.6
2.4
2.2
2.0
1.8
1.6
0
1000
2000
3000
4000
5000
Completed optimizer updates